

exercice2

KNN PRACTICE

```
library(tidyverse)

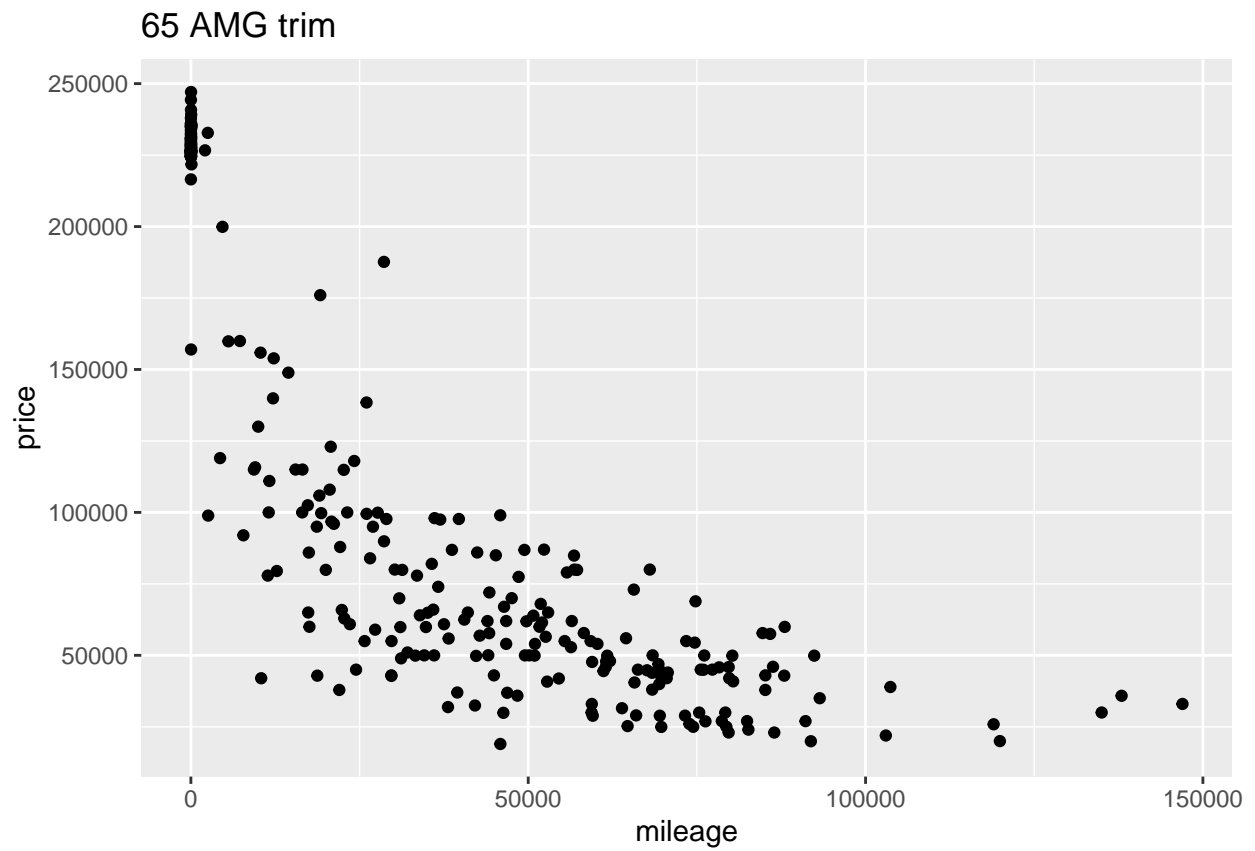
## -- Attaching packages ----- tidyverse 1.3.0 --
## v ggplot2 3.2.1    v purrr  0.3.3
## v tibble  2.1.3    v dplyr  0.8.3
## v tidyr   1.0.0    v stringr 1.4.0
## v readr   1.3.1    v forcats 0.4.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()


library(ggplot2)
sclass <- read.csv("~/Documents/R/SDS 323/SDS323-master/data/sclass.csv")

# trim subsets
sclass350 <- dplyr::select(subset(sclass, trim == "350"), mileage, price)
sclass65AMG <- dplyr::select(subset(sclass, trim == "65 AMG"), mileage, price)

# plot price vs mileage for each trim
ggplot(data = sclass65AMG) + geom_point(aes(x = mileage, y = price)) + labs(title = "65 AMG trim")
```



```
ggplot(data = sclass350) + geom_point(aes(x = mileage, y = price)) + labs(title = "350 trim")
```



Looking at the two plots, the data seems to indicate different relationships between mileage and price for each S Class. Hence using a separate KNN regression model for each S Class seems to be a good idea.

```
library(FNN)
library(mosaic)
```

```
## Loading required package: lattice
## Loading required package: ggformula
## Loading required package: ggstance
##
## Attaching package: 'ggstance'
##
## The following objects are masked from 'package:ggplot2':
##
##   geom_errorbarh, GeomErrorbarh
##
## New to ggformula? Try the tutorials:
##   learnr::run_tutorial("introduction", package = "ggformula")
##   learnr::run_tutorial("refining", package = "ggformula")
## Loading required package: mosaicData
## Loading required package: Matrix
##
## Attaching package: 'Matrix'
##
## The following objects are masked from 'package:tidyr':
##
```

```

##      expand, pack, unpack
## Registered S3 method overwritten by 'mosaic':
##      method                      from
##      fortify.SpatialPolygonsDataFrame ggplot2
##
## The 'mosaic' package masks several functions from core packages in order to add
## additional features. The original behavior of these functions should not be affected by this.
##
## Note: If you use the Matrix package, be sure to load it BEFORE loading mosaic.
##
## Attaching package: 'mosaic'
##
## The following object is masked from 'package:Matrix':
##
##      mean
##
## The following objects are masked from 'package:dplyr':
##
##      count, do, tally
##
## The following object is masked from 'package:purrr':
##
##      cross
##
## The following object is masked from 'package:ggplot2':
##
##      stat
##
## The following objects are masked from 'package:stats':
##
##      binom.test, cor, cor.test, cov, fivenum, IQR, median, prop.test,
##      quantile, sd, t.test, var
##
## The following objects are masked from 'package:base':
##
##      max, mean, min, prod, range, sample, sum
library(doMC)

## Loading required package: foreach
##
## Attaching package: 'foreach'
##
## The following objects are masked from 'package:purrr':
##
##      accumulate, when
##
## Loading required package: iterators
## Loading required package: parallel
library(boot)

##
## Attaching package: 'boot'
##
## The following object is masked from 'package:mosaic':
##

```

```

##      logit
## The following object is masked from 'package:lattice':
##
##      melanoma

options(`mosaic:parallelMessage` = FALSE)
set.seed(100)
set.rseed(100)

rmse <- function(y, yhat) {
  sqrt( mean( (y - yhat)^2 ) )
}

# KNN regression for S Class 350
X <- dplyr::select(sclass350, -price)
y <- sclass350$price
n <- nrow(sclass350)
train_ind <- n * 0.8


k_grid <- seq(1,80,by=1)
err_grid <- foreach(k = k_grid, .combine = 'c') %do% {
  out = do(500)*{
    # test/train split
    train_ind <- sample.int(nrow(sclass350), 0.8*nrow(sclass350))
    X_train <- data.frame(X[train_ind,])
    X_test <- data.frame(X[-train_ind,])
    y_train <- y[train_ind]
    y_test <- y[-train_ind]

    # scale train and test feature by the sd of train features
    scale_factors <- apply(X_train, 2, sd, na.rm = TRUE)
    X_train_sc <- scale(X_train, scale = scale_factors)
    X_test_sc <- scale(X_test, scale = scale_factors)

    model <- knn.reg(train = X_train_sc, test = X_test_sc, y = y_train, k = k)

    rmse(y_test, model$pred)
  }
  mean(out$result)
}
plot(err_grid)

```



```
# index of optimal K and RMSE
which.min(err_grid)

## [1] 37

min(err_grid)

## [1] 10520.23


# plot optimal K
scale_factors <- apply(X, 2, sd, na.rm = TRUE)
X_test <- seq(0, 200000, length.out = 1000)
X_train_sc <- scale(X_train, scale = scale_factors)
X_test_sc <- scale(X_test, scale = scale_factors)
best <- knn.reg(train = X_train_sc, test = X_test_sc, y = y, k = which.min(err_grid))

# quadratic model for comparisson
lm <- glm(price ~ poly(x = mileage, degree = 2), data = sclass350)
lm.cv <- cv.glm(data = sclass350, lm)

# rmse of quadratic model
sqrt(lm.cv$delta[1])

## [1] 10238.95

df1 <- data.frame(X_test, best$pred)
df2 <- data.frame(X_test, predict(lm, newdata = data.frame(mileage = X_test)))
plot(price ~ mileage, data = sclass350, col = "black")
lines(df1, col = "blue")
lines(df2, col = "red")
```



The

optimal value for K is 37, which happens to be odd, and has an RMSE of 10520.23.

```
library(doMC)
set.seed(100)
set.rseed(100)
# KNN regression for S Class 65 AMG

X <- dplyr::select(sclass65AMG, -price)
y <- sclass65AMG$price

n <- nrow(sclass65AMG)
train_n <- n*0.8


k_grid <- seq(1,80,by=1)
err_grid <- foreach(k = k_grid, .combine = 'c') %do% {
  out = do(500)*{
    # make test/train split
    train_ind <- sample.int(n, train_n)
    X_train <- data.frame(X[train_ind,])
    X_test <- data.frame(X[-train_ind,])
    y_train <- y[train_ind]
    y_test <- y[-train_ind]

    # scale train and test feature by the sd of train features
    scale_factors <- apply(X_train, 2, sd, na.rm = TRUE)
    X_train_sc <- scale(X_train, scale = scale_factors)
    X_test_sc <- scale(X_test, scale = scale_factors)

    model <- knn.reg(train = X_train_sc, test = X_test_sc, y = y_train, k = k)

    rmse(y_test, model$pred)
  }
  mean(out$result)
}
```

```
}
plot(err_grid)
```



```
# optimal K and RMSE
which.min(err_grid)
```

```
## [1] 36
```

```
min(err_grid)
```

```
## [1] 26012.75
```

```
# optimal odd K and RMSE
odd <- err_grid[c(TRUE, FALSE)]
2*which.min(odd)-1
```

```
## [1] 35
```

```
min(odd)
```

```
## [1] 26087.95
```

```
# RMSE difference
min(err_grid) - min(odd)
```

```
## [1] -75.20318
```

```
lm <- glm(price ~ poly(x = mileage, degree = 5), data = sclass65AMG)
lm.cv <- cv.glm(sclass65AMG, lm)
sqrt(lm.cv$delta[1])
```

```
## [1] 21194.57
```


```
X_test <- seq(0, 150000, length.out = 1000)
scale_factors <- apply(X, 2, sd, na.rm = TRUE)
X_train_sc <- scale(X, scale = scale_factors)
X_test_sc <- scale(X_test, scale = scale_factors)
```



```
best <- knn.reg(X_train_sc, X_test_sc, y = y, k = which.min(err_grid))

df1 <- data.frame(X_test, best$pred)
df2 <- data.frame(X_test, predict(lm, newdata = data.frame(mileage = X_test)))

plot(price ~ mileage, data = sclass65AMG, col = "red")
lines(df1, col = "blue")
lines(df2, col = "black")
```



The model for 350 trim has the larger optimal K, this is likely because the data set is larger than that of 65 AMG trim. Therefore the model can average more data points without losing accuracy. The optimal K value is 36 and the optimal odd K is 35, the RMSE difference between the two is ~75. Because the two values for K are neighbors and the graph doesn't seem volatile, I think it is safe to use K = 36 as the optimal value for K. It is interesting to note that for both trims a simple polynomial model outperformed the KNN regression model, and is substantially easier to implement.

SARATOGA HOUSE PRICES

```
set.seed(100)
set.rseed(100)
library(mosaic)
library(boot)
library(lm.beta)
data(SaratogaHouses)

summary(SaratogaHouses)
```

##	price	lotSize	age	landValue
## Min.	: 5000	Min. : 0.0000	Min. : 0.00	Min. : 200
## 1st Qu.	: 145000	1st Qu.: 0.1700	1st Qu.: 13.00	1st Qu.: 15100
## Median	: 189900	Median : 0.3700	Median : 19.00	Median : 25000
## Mean	: 211967	Mean : 0.5002	Mean : 27.92	Mean : 34557
## 3rd Qu.	: 259000	3rd Qu.: 0.5400	3rd Qu.: 34.00	3rd Qu.: 40200

```

## Max. :775000 Max. :12.2000 Max. :225.00 Max. :412600
## livingArea pctCollege bedrooms fireplaces bathrooms
## Min. : 616 Min. :20.00 Min. :1.000 Min. :0.0000 Min. :0.0
## 1st Qu.:1300 1st Qu.:52.00 1st Qu.:3.000 1st Qu.:0.0000 1st Qu.:1.5
## Median :1634 Median :57.00 Median :3.000 Median :1.0000 Median :2.0
## Mean :1755 Mean :55.57 Mean :3.155 Mean :0.6019 Mean :1.9
## 3rd Qu.:2138 3rd Qu.:64.00 3rd Qu.:4.000 3rd Qu.:1.0000 3rd Qu.:2.5
## Max. :5228 Max. :82.00 Max. :7.000 Max. :4.0000 Max. :4.5
## rooms heating fuel
## Min. : 2.000 hot air :1121 gas :1197
## 1st Qu.: 5.000 hot water/steam: 302 electric: 315
## Median : 7.000 electric : 305 oil : 216
## Mean : 7.042
## 3rd Qu.: 8.250
## Max. :12.000
## sewer waterfront newConstruction centralAir
## septic : 503 Yes: 15 Yes: 81 Yes: 635
## public/commercial:1213 No :1713 No :1647 No :1093
## none : 12
##
##
##
# baseline medium model with 11 main effects
lm_medium <- glm(price ~ lotSize + age + livingArea + pctCollege + bedrooms +
  fireplaces + bathrooms + rooms + heating + fuel + centralAir, data=SaratogaHouses)
summary(lm_medium)

##
## Call:
## glm(formula = price ~ lotSize + age + livingArea + pctCollege +
## bedrooms + fireplaces + bathrooms + rooms + heating + fuel +
## centralAir, data = SaratogaHouses)
##
## Deviance Residuals:
## Min 1Q Median 3Q Max
## -232296 -40021 -7679 28919 527748
##
## Coefficients:
## Estimate Std. Error t value Pr(>|t|)
## (Intercept) 28627.732 12224.396 2.342 0.019302 *
## lotSize 9350.452 2421.118 3.862 0.000117 ***
## age 47.547 65.149 0.730 0.465600
## livingArea 91.870 5.033 18.253 < 2e-16 ***
## pctCollege 296.508 165.531 1.791 0.073428 .
## bedrooms -15630.719 2885.084 -5.418 6.89e-08 ***
## fireplaces 985.061 3385.478 0.291 0.771112
## bathrooms 22006.971 3821.764 5.758 1.00e-08 ***
## rooms 3259.119 1093.631 2.980 0.002922 **
## heatinghot water/steam -9429.795 4738.934 -1.990 0.046765 *
## heatingelectric -3609.986 14009.898 -0.258 0.796689
## fuelelectric -12094.122 13792.538 -0.877 0.380686
## fueloil -8873.140 5395.649 -1.644 0.100257
## centralAirNo -17112.819 3922.489 -4.363 1.36e-05 ***
## ---

```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 4393644476)
##
##      Null deviance: 1.6736e+13  on 1727  degrees of freedom
## Residual deviance: 7.5307e+12  on 1714  degrees of freedom
## AIC: 43287
##
## Number of Fisher Scoring iterations: 2

# K=10 CV
lm_medium.cv <- cv.glm(data = SaratogaHouses, lm_medium, K = 10)

# RMSE of the medium model
sqrt(lm_medium.cv$delta[1])

## [1] 66672.84

# my model
lm1 <- glm(price ~ lotSize + livingArea + pctCollege + bedrooms + bathrooms + rooms + centralAir + waterfront, data = SaratogaHouses)
lm1.cv <- cv.glm(data = SaratogaHouses, lm1, K = 10)

# RMSE of my model
sqrt(lm1.cv$delta[1])

## [1] 64900.97

# RMSE improvement on the medium model
sqrt(lm_medium.cv$delta[1]) - sqrt(lm1.cv$delta[1])

## [1] 1771.873

lm <- lm(price ~ lotSize + livingArea + pctCollege + bedrooms + bathrooms + rooms + centralAir + waterfront, data = SaratogaHouses)
lm.beta(lm)

##
## Call:
## lm(formula = price ~ lotSize + livingArea + pctCollege + bedrooms +
##      bathrooms + rooms + centralAir + waterfront, data = SaratogaHouses)
##
## Standardized Coefficients:
##      (Intercept)      lotSize    livingArea    pctCollege    bedrooms    bathrooms
##      0.00000000    0.05741375    0.56569250    0.04993521   -0.10582911    0.14661605
##           rooms centralAirNo waterfrontNo
##      0.08359965   -0.10446242   -0.15881708
```

Adding the waterfront feature and removing age, fireplaces, heating, fuel, and centralAir from the medium model decreases the RMSE by about \$1408. I added the waterfront feature because in my experience, waterfront properties are more expensive than their more land-locked partners. I chose to remove many of the features from the medium model because their coefficients were not statistically significantly non-zero. These features may not have been statistically significant, because they do not provide additional information for the model, or because of collinearity. In that case adding the interaction terms would help resolve this issue. But for this assignment, simply removing those features led to a significant improvement in the model's accuracy. Using the standardized coefficients we can see that livingArea, waterfront, and bathrooms have a large impact on the predictions of our model.

```
library(FNN)
library(caret)
```

```

##
## Attaching package: 'caret'

## The following object is masked from 'package:mosaic':
##
##      dotPlot

## The following object is masked from 'package:purrr':
##
##      lift

library(mosaic)
library(doMC) # parallel computing
set.seed(100)
set.rseed(100)
data(SaratogaHouses)

rmse = function(y, yhat) {
  sqrt( mean( (y - yhat)^2 ) )
}

# one-hot encode the categorical variables
dmy <- dummyVars("~.", data = SaratogaHouses)

# create data frame with new variables
data <- data.frame(predict(dmy, newdata=SaratogaHouses))

# dont use price
X <- data[,2:ncol(data)]
y <- data$price
n <- nrow(data)
train_n <- n * 0.8

# knn regression
k_grid <- seq(1,50,by=1)
k7 <- rep(0, 500)
for(k in k_grid) {
  err <- rep(0, 500)

  for(i in 1:length(err)) {
    train_ind <- sample.int(n, train_n)
    X_train <- X[train_ind,]
    X_test <- X[-train_ind,]
    y_train <- y[train_ind]
    y_test <- y[-train_ind]

    # scale train and test feature by the sd of train features
    scale_factors <- apply(X_train, 2, sd, na.rm = TRUE)
    X_train_sc <- scale(X_train, scale = scale_factors)
    X_test_sc <- scale(X_test, scale = scale_factors)

    model <- knn.reg(train = X_train_sc, test = X_test_sc, y = y_train, k = k)

    err[i] <- rmse(y_test, model$pred)
    if(k == 7) k7[i] <- err[i]
  }
}


```

```

}
k_grid[k] <- mean(err)
}

# plot the test average RMSE
plot(k_grid)

```



```

# find the optimal K value
which.min(err)

```

```
## [1] 127
```

```
min(err)
```

```
## [1] 54498.67
```

```

# improvement on the hand crafted model
sqrt(lm1.cv$delta[1]) - min(err)

```

```
## [1] 10402.3
```

```

# t test
t.test(k7)

```

```

##
## One Sample t-test
##
## data: k7
## t = 288.41, df = 499, p-value < 2.2e-16
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
## 63404.86 64274.65
## sample estimates:
## mean of x
## 63839.75

```

The optimal value for K is 7 with an RMSE of \$63772.17, which is an improvement of about \$1128.80 over my hand crafted model. The 95% confidence interval for the RMSE is (63404, 64274). -REPORT- Our model has an expected RMSE of \$63772, meaning it is likely that our predictions for house price will regularly be off by about \$60000. This is likely to have a large impact on the tax rate we charge each household, however I do not think it will have a large impact on the tax bracket of each house. Because of the large in-accruary of our model, I think it is important to factor in past evaluations of the house when deciding what tax rate to charge each household.

PREDICTING WHEN ARTICLES GO VIRAL

```
library(glmnet)

## Loaded glmnet 3.0-2


library(tidyverse)
library(boot)
library(caret)
set.seed(100)

news <- read.csv("~/Documents/R/SDS 323/SDS323-master/data/online_news.csv")
news <- dplyr::select(news, -url)

# calculate our null prediction
null_prediction <- sum(news$shares>1400)/nrow(news)

# seperate the data into features and response
X <- model.matrix(shares ~. -1, data=news)
y <- news$shares

# find the optimal lambda for our data set
cv.lasso1 <- cv.glmnet(X, y, alpha = 1, family = "gaussian", nfolds = 10)
plot(cv.lasso1)
```



```

# K=10 CV
K <- 10
fold_id <- rep_len(1:K, nrow(X))
fold_id <- sample(fold_id)

acc<- rep(0, K)
tp<- rep(0, K)
fp<- rep(0, K)
for(i in 1:K) {
  train_set <- which(fold_id != i)
  X_train <- X[train_set,]
  X_test <- X[-train_set,]
  y_train <- y[train_set]
  y_test <- y[-train_set]
  y_test <- ifelse(y_test > 1400, 1, 0)

  l1 <- glmnet(X_train, y_train, alpha = 1, family = "gaussian", lambda = cv.lasso1$lambda.min)
  y_hat <- predict(l1, X_test)
  y_hat <- ifelse(y_hat > 1400, 1, 0)
  cm <- confusionMatrix(factor(y_hat), factor(y_test))$table
  cm
  acc[i] <- (cm[1] + cm[4])/sum(cm)
  tp[i] <- cm[4]/(cm[3] + cm[4])
  fp[i] <- cm[2]/(cm[1] + cm[2])
}
# accuracy
mean(acc)

```

```
## [1] 0.49662
```

```

# true positive rate
mean(tp)

```

```
## [1] 0.9960667
```

```

# false positive rate
mean(fp)

```

```
## [1] 0.9898802
```

Training our model before thresholding y leads to a model that is less accurate than the no information rate, meaning it is less accurate than simply using the sample proportion of viral articles to predict viral status.


Now we perform the same analysis, but we threshold y first.

```

set.seed(100)
viral <- news$shares
viral <- ifelse(viral > 1400, 1, 0)
y <- viral

# find the optimal lambda for the data set
X <- model.matrix(shares ~. -1, data=news)
cv.lasso2 <- cv.glmnet(X, y, alpha = 1, family = "binomial", nfolds = 10)
plot(cv.lasso2)

```



```
# K=10 CV
K <- 10
l2_err <- rep(0, K)
fold_id <- rep_len(1:K, nrow(X))
fold_id <- sample(fold_id)

for(i in 1:K) {
  train_set <- which(fold_id != i)
  X_train <- X[train_set,]
  X_test <- X[-train_set,]
  y_train <- y[train_set]
  y_test <- y[-train_set]

  l2 <- glmnet(X_train, y_train, alpha = 1, family = "binomial", lambda = cv.lasso2$lambda.min)
  y_hat <- predict(l2, X_test)
  y_hat <- sapply(y_hat, function(x){ifelse(x > 0.5, 1, 0)})

  l2_err[i] <- mean(y_hat == y_test)
}
mean(l2_err)
```

```
## [1] 0.5907576
```

```
# logistic model using every variable
X <- dplyr::select(news, -shares)

# K fold CV
acc<- rep(0, K)
tp<- rep(0, K)
fp<- rep(0, K)
```



```
fold_id <- rep_len(1:K, nrow(X))
fold_id <- sample(fold_id)
for(k in 1:K) {
  train_set <- which(fold_id != k)
  X_train <- X[train_set,]
  X_test <- X[-train_set,]
  y_train <- y[train_set]
  y_test <- y[-train_set]

  model <- lm(shares ~., data = data.frame(X_train, shares = y_train))
  y_hat <- predict(model, newdata = X_test)
  y_hat <- ifelse(y_hat > 0.5, 1, 0)
  cm <- confusionMatrix(factor(y_hat), factor(y_test))$table
  acc[i] <- (cm[1] + cm[4])/sum(cm)
  tp[i] <- cm[4]/(cm[4] + cm[3])
  fp[i] <- cm[2]/(cm[1] + cm[2])
}

## Warning in predict.lm(model, newdata = X_test): prediction from a rank-deficient
## fit may be misleading

## Warning in predict.lm(model, newdata = X_test): prediction from a rank-deficient
## fit may be misleading

## Warning in predict.lm(model, newdata = X_test): prediction from a rank-deficient
## fit may be misleading

## Warning in predict.lm(model, newdata = X_test): prediction from a rank-deficient
## fit may be misleading

## Warning in predict.lm(model, newdata = X_test): prediction from a rank-deficient
## fit may be misleading

## Warning in predict.lm(model, newdata = X_test): prediction from a rank-deficient
## fit may be misleading

## Warning in predict.lm(model, newdata = X_test): prediction from a rank-deficient
## fit may be misleading

## Warning in predict.lm(model, newdata = X_test): prediction from a rank-deficient
## fit may be misleading

## Warning in predict.lm(model, newdata = X_test): prediction from a rank-deficient
## fit may be misleading

# accuracy
mean(acc)

## [1] 0.06291625

# true positive rate
mean(tp)
```

```
## [1] 0.06321897
```

```
# false positive rate  
mean(fp)
```

```
## [1] 0.03738648
```

KNN regression code.

```
# library(FNN)  
# library(matrixStats)  
#  
# colScale <- function(x,  
#   center = TRUE,  
#   scale = TRUE,  
#   add_attr = TRUE,  
#   rows = NULL,  
#   cols = NULL) {  
#  
#   if (!is.null(rows) && !is.null(cols)) {  
#     x <- x[rows, cols, drop = FALSE]  
#   } else if (!is.null(rows)) {  
#     x <- x[rows, , drop = FALSE]  
#   } else if (!is.null(cols)) {  
#     x <- x[, cols, drop = FALSE]  
#   }  
#  
#   #####  
#   # Get the column means  
#   #####  
#   cm = colMeans(x, na.rm = TRUE)  
#   #####  
#   # Get the column sd  
#   #####  
#   if (scale) {  
#     csd = colSds(x, center = cm)  
#   } else {  
#     # just divide by 1 if not  
#     csd = rep(1, length = length(cm))  
#   }  
#   if (!center) {  
#     # just subtract 0  
#     cm = rep(0, length = length(cm))  
#   }  
#   x = t( (t(x) - cm) / csd )  
#   if (add_attr) {  
#     if (center) {  
#       attr(x, "scaled:center") <- cm  
#     }  
#     if (scale) {  
#       attr(x, "scaled:scale") <- csd  
#     }  
#   }  
#   return(x)  
# }  
#
```

```

# # KNN regression with 1se optimal lambda coefficients
# X <- model.matrix(shares ~. -1, data = news)
# l2 <- glmnet(X, y, alpha = 1, family = "binomial", lambda = cv.lasso2$lambda.1se)
#
# # get non zero coefficients
# lasso_coefs <- rownames(coef(l2))[coef(l2)[,1] == 0]
#
#
# X <- dplyr::select(news, -c(shares, lasso_coefs))
# sample_ind <- sample.int(nrow(X), 1000)
# sX <- X[sample_ind,]
# sy <- y[sample_ind]
#
#
# n <- nrow(X)
# train_n <- n*0.8
# k_grid <- seq(40,100, by = 1)
# for(k in k_grid) {
#   err <- rep(0, 5)
#   for(i in 1:length(err)) {
#     train_ind <- sample.int(n, train_n)
#     X_train <- X[train_ind,]
#     X_test <- X[-train_ind,]
#     y_train <- y[train_ind]
#     y_test <- y[-train_ind]
#
#     #scale_factors <- apply(X_train, 2, sd, na.rm = TRUE)
#     X_train_sc <- colScale(X_train, center = FALSE, scale = TRUE)
#     X_test_sc <- colScale(X_test, center = FALSE, scale = TRUE)
#     #X_train_sc <- scale(X_train, scale = scale_factors)
#     #X_test_sc <- scale(X_test, scale = scale_factors)
#
#     model <- knn.reg(X_train_sc, X_test_sc, y_train, k)
#     y_hat <- model$pred
#     y_hat <- ifelse(model$pred > 0.5, 1, 0)
#
#     err[i] <- mean(y_test == y_hat)
#     print(i)
#   }
#
#   k_grid[k] <- mean(err)
# }
#
# library(ggplot2)
# df <- data.frame(x = seq(1,100,by = 1), y = k_grid)
# ggplot(data = df) + xlim(40,100) + ylim(0.6,0.65) + geom_point(aes(x = x, y = y), data= df)
# plot(k_grid)
#
# which.max(k_grid)
# View(k_grid)
# k_grid[which.max(k_grid)]

```

Running the regression before thresholding y lead to 49.7% accuracy, slightly above the null prediction of 49.3%, which is the percentage of “viral” articles in the data set. Thresholding y before running the regression

increased the accuracy to 59% an improvement of about 10% over our other model. We can see why this may be by looking at the MSE vs $\log(\lambda)$ plots for each model. In the first plot, we can see there is minimal change in MSE for varying values for $\log(\lambda)$, so much so that the difference in MSE corresponding to the min and max values for $\log(\lambda)$ is negligible. More importantly, the confidence intervals are extremely large and thus it is possible the “optimal” value of λ is not actually optimal. These observations are in stark contrast to those of the plot for model 2 (thresholding y before regression). In these plots, the confidence intervals are smaller and the plots look more similar to what we should expect to see. I also ran a logistic regression model using every parameter (excluding URL) and that saw an improvement to 63.7% accuracy. I also attempted to run a KNN regression, but with my computer’s limited resources and the large data set (about 40000*36 data points), I could not complete the regression in any reasonable amount of time. I only used the non-zero coefficients from the 1 se λ calculated from my previous lasso regression, I knew I may be throwing away information, but I hoped cutting out ~10 parameters would drastically increase the speed of training the model. This seemed to work, but it certainly wasn’t enough to make the computation practical in a reasonable about of time. My next idea was to randomly sample a smaller portion of the data set and train the model on that data set. However this was unreliable, as the optimal K would change greatly for different sample. I could fix this by increasing the sample size, but then that defeats the purpose because I run into the original problem. Finally I decided to implement a faster scale method, because from my testing that seemed to be one of the major resource drains in training the model (that along with calculating the distances for each point). I found a column scale function on the internet that seemed to be faster and implemented it. It greatly increased the speed of the training algorithm. I trained the model over the entire data set 5 times for each K between 40 and 100 (the previous models made me confident that the optimal K was at least greater than 40), this took about 2 hours and was about 63.3% accurate for the optimal K . But to gain confidence in the model I would have to train it more than 5 times (most likely several hundred) for each value of K , which was still completely impractical. So I decided to scrap the idea of running a KNN regression model, but I left my code if you wanted to take a look.